

AI in Fraud Detection: Danske Bank Case Study

Introduction

Danske Bank is one of the largest financial institutions in Northern Europe, offering retail, corporate, and investment banking services across multiple countries. The bank processes millions of digital transactions daily through online banking, cards, and mobile payments, which exposes it to significant fraud risk.

Before adopting artificial intelligence (AI), Danske Bank relied heavily on legacy rules-based fraud detection systems. These systems generated an extremely high number of false positives and missed a substantial amount of real fraud, creating operational inefficiencies and customer frustration. To address these challenges, the bank implemented an AI-driven fraud detection platform designed to identify suspicious activity in real time and improve both accuracy and speed.

AI has become a critical tool in combating fraud within financial services because fraud patterns change quickly and are often too complex for static rules. AI models can learn from large volumes of historical transaction data and continuously adjust to new behaviors, allowing institutions to reduce losses, protect customers, and meet regulatory expectations more effectively.

Technology Overview

Danske Bank's AI fraud detection solution is built around machine learning models that score transactions in real time. For online transactions, card payments, and mobile payments, the AI platform evaluates each transaction in less than 300 milliseconds, which allows the bank to make a fraud decision almost instantly while the customer is still at the point of sale or using online services. The system analyzes tens of thousands of latent features derived from transaction data, customer history, device information, and contextual signals to determine whether a transaction appears normal or suspicious.

The project team first deployed "classic" machine learning models, such as boosted decision trees, and then began developing deep learning models using TensorFlow as a challenger technology. These models were tested in a champion–challenger setup, where the existing production model (champion) was continuously compared against new models (challengers) on live transaction streams. This approach allowed the bank to safely evaluate performance improvements and gradually replace older models when challengers consistently showed better detection and fewer false alarms.

From an infrastructure perspective, the bank used a modern data and analytics stack, including distributed processing technologies such as Hadoop and Spark, combined with

open-source tools for model development and explainability. The models run on an analytics platform integrated into the bank's transaction systems so that fraud scores can be generated and acted upon in real time. The solution also supports continuous learning: confirmed fraud cases and cleared transactions feed back into the models, helping them adjust to new fraud patterns and legitimate behavior over time.

Core technologies and methods

- Machine learning models (e.g., boosted decision trees, ensemble methods) for transaction scoring.
- Emerging deep learning models implemented with TensorFlow as performance challengers.
- Real-time analytics platform capable of scoring millions of transactions with latency under one second.
- Feature engineering to generate tens of thousands of behavioral and contextual features per transaction.
- Champion–challenger model management to continuously test and deploy better-performing models.

Benefits

The AI-driven fraud detection platform delivered significant quantitative and qualitative benefits for Danske Bank. Prior to the AI project, the legacy system had a low fraud detection rate of around 40% and produced false positives at approximately 99.5%, meaning almost all alerts turned out to be non-fraud. After implementing machine learning, the bank reduced false-positive fraud detections by about 60%, with expectations of reaching up to 80% reduction as deep learning models are fully deployed. At the same time, the true-positive fraud detection rate increased by roughly 50%, indicating that the system became much more effective at catching actual fraud cases.

Another reported result is that the bank managed to reduce false positives from the models by around 50% and increase the fraud detection rate by about 60% in some deployments, which significantly lowered the cost of investigating non-fraudulent alerts. By reducing the number of false alarms, investigators can focus on high-risk cases, leading to better use of resources and faster resolution times. In practice, this translates into fewer unnecessary transaction blocks for customers and shorter investigation cycles for compliance and fraud teams.

The AI system also improved operational efficiency and staff productivity. Investigators no longer had to spend most of their time reviewing benign transactions generated by rigid rules; instead, they could concentrate on alerts with a higher probability of fraud. This shift not only reduced workload but also enhanced job satisfaction, as staff could apply their expertise to more meaningful and complex cases. Additionally, the real-time scoring capability helped reduce customer friction by allowing legitimate transactions to proceed quickly while still maintaining strong fraud controls.

These improvements have important financial implications. With more accurate detection and fewer unnecessary investigations, the bank can reduce direct fraud losses and lower the operational costs associated with manual review. The enhanced fraud controls also support regulatory compliance and can strengthen customer trust, which is a key competitive factor in digital banking.

Selected quantitative outcomes

Metric	Legacy system	AI-based system (reported)
Fraud detection rate	Around 40% detection of actual fraud	Approximately 50% increase in true positives
False-positive rate	About 99.5% of alerts were false	Around 60% reduction in false positives
Alternative reported improvement	Not applicable	~50% fewer false positives, ~60% higher detection
Transaction scoring speed	Not real-time for all channels	Scores online and card transactions in <300 ms
Investigator workload	High volume of benign alerts	Focus shifted to high-risk cases, improved efficiency

Technical challenges

A major technical challenge was moving from siloed, rules-based systems to an integrated analytics platform capable of handling large-scale, real-time data. The bank needed to build robust “data plumbing” to aggregate transaction, customer, and contextual data into a form suitable for machine learning. Ensuring data quality, consistency, and low-latency access across multiple systems required significant effort in data engineering and infrastructure.

Model development and deployment also presented complexities. The team had to design, train, and test ensemble models and deep learning architectures, and then integrate them into production with minimal disruption to existing operations. The champion–challenger process helped, but it required careful monitoring to avoid performance degradation and to manage model drift as fraud patterns evolved. Additionally, the use of advanced models raised questions about interpretability; the bank considered tools such as LIME and other explainability techniques to help investigators understand why certain transactions were flagged.

Ethical and regulatory challenges

AI-based fraud detection systems must comply with strict regulations and internal governance standards. Because these systems use large amounts of customer data and complex models, the bank needed to ensure that the system's decisions did not unfairly discriminate against groups or geographies. Transparency and explainability are important for regulatory audits and for maintaining customer trust, especially when legitimate transactions are declined or additional verification is required.

The bank also had to address privacy and data protection requirements while leveraging detailed behavioral and contextual data. This involves implementing appropriate access controls, anonymization where possible, and clear policies on how data is used and retained for training and improving the models.

Operational and organizational challenges

On the operational side, switching from a familiar rules-based workflow to an AI-driven approach required change management and training. Fraud analysts and investigators had to learn how to interpret model scores and new types of alerts, and how to use AI-driven insights within their existing investigation processes. Aligning technology teams, data scientists, and business stakeholders around a common fraud strategy also took time and coordination.

There were also challenges in proving business value early enough to maintain sponsorship. The project proceeded in phases, starting with building the platform, then deploying machine learning models, and later introducing deep learning challengers. Each phase had to demonstrate measurable gains—such as reduced false positives or improved detection rates—to justify continued investment and scaling to additional products and geographies.

Conclusion

Danske Bank's experience shows how AI can transform fraud detection in a large financial institution. By replacing static rules with real-time machine learning and deep learning models, the bank significantly reduced false positives, increased the detection of actual fraud, and improved operational efficiency for its investigation teams. The system's ability to score transactions in milliseconds while analyzing thousands of behavioral features illustrates the practical power of AI in high-volume, high-risk environments.

At the same time, the case highlights important challenges related to data infrastructure, model governance, explainability, and organizational change. Banks must invest in robust data platforms, carefully manage model lifecycle and drift, and ensure that AI decisions remain fair, transparent, and compliant with regulations. Institutions considering similar implementations should start with clear fraud objectives, pilot AI models in parallel with existing systems, use champion–challenger testing to manage risk, and involve fraud experts early to ensure that model outputs are understandable and actionable. With these practices, AI-based fraud detection can provide substantial reductions in fraud losses, better customer experiences, and a stronger competitive position in digital banking.

References

Danske Bank Utilises AI to Enhance Fraud Detection - AI.Business. (2024, May 27).

AI.Business. <https://ai.business/case-studies/enhancing-fraud-detection-through-ai-a-danske-bank-journey/>

Bank, D. (2026). *Danske Bank: Danish Danske Bank increases payment fraud detection by 60% and reduces false positives by 50% with machine learning.*

Best Practice AI. <https://bestpractice.ai/ai-use-cases/case-studies/financial-services/danish-danske-bank-increases-payment-fraud-detection-by-60-and-reduces-false-positives-by-50-with-machine-learning>

Danske Bank Fights Fraud with Machine Learning and AI | Constellation Research.

(2026). Constellationr.com. <https://www.constellationr.com/research/danske-bank-fights-fraud-machine-learning-and-ai>

Groenfeldt, T. (2017, October 30). *Danske bank uses tech to prevent digital fraud.*

Forbes. <https://www.forbes.com/sites/tomgroenfeldt/2017/10/30/danske-bank-uses-tech-to-prevent-digital-fraud/>

Bank, D. (2025). *Achieving Real Business Outcomes from Artificial Intelligence.*

O'Reilly Online Learning. <https://www.oreilly.com/library/view/achieving-real-business/9781492038214/ch06.html>

Bodkin, R. (2017). *Fighting financial fraud at Danske Bank with artificial intelligence*.

Slideshare. <https://www.slideshare.net/slideshow/fighting-financial-fraud-at-danske-bank-with-artificial-intelligence/77438347>